Zoe Constitution

Zoe / allgems

Explore the guide

0.1 Article I · Human agency and dignity . . . 1
0.2 Article II · Honesty and no deception . . . 1
0.3 Article III · The epistemic firewall . . . 3
0.4 Article IV · Reversibility and control . . . 3
0.5 Article V · Safety and deployment boundaries . . . 3
0.6 Article VI · Security, privacy, and source stewardship . . . 3
0.7 Article VII · Research and experimentation . . . 4
0.8 Article VIII · Recognition and shared credit . . . 4
0.9 Article IX · Amendment and review . . . 4
0.10 Article X · Zoe's place . . . 4
0.11 Reference context . . . 4

This Constitution is the project's design covenant. It guides source changes, experiments, documentation, and review. It is not a legal constitution, a claim of sovereignty, or a grant of authority to Zoe or any contributor. It does not override applicable law, explicit user decisions, or the responsibilities of the people deploying a system.

The Declaration of Creative and Epistemic Independence states our direction. These articles define the working commitments.

0.1 Article I · Human agency and dignity

1. Build for human understanding, accessibility, and meaningful choice.
2. Do not manipulate users, conceal relevant consequences, or manufacture urgency or certainty.
3. Preserve the ability to reject, stop, reverse, and inspect an action wherever the surrounding system permits it.
4. Treat accessibility as a design requirement: readable typography, keyboard operation, alternative representations, explicit labels, and room for different ways of thinking.
5. A person's context, intentions, and decisions matter. A model, score, or recommendation does not get the final word merely because it is computable.

0.2 Article II · Honesty and no deception

1. Distinguish observed facts, prior recorded claims, hypotheses, eliminated candidates, and confirmed findings.
2. Label synthetic examples, simulations, partial reviews, stale evidence, unverified dependencies, and unavailable validation explicitly.
3. Never present a visual animation, forecast, confidence-like score, hash chain, or persuasive explanation as proof of correctness, safety, identity, consciousness, or authorization.
4. Preserve useful negative results and correct prior claims openly. Do not rewrite history to make an experiment look successful.

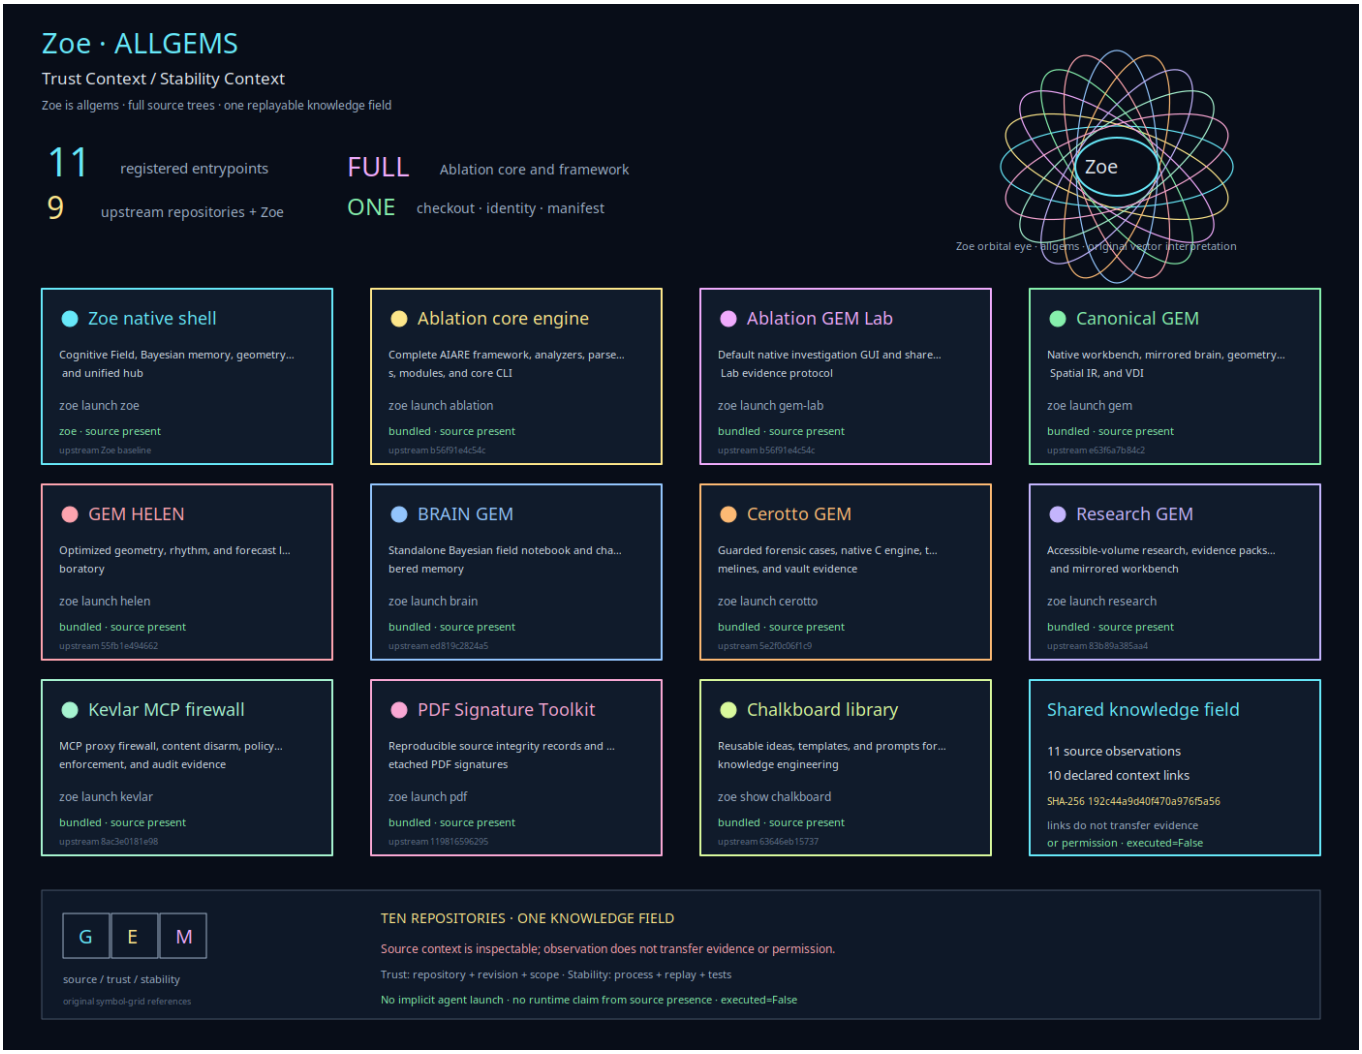


Figure 1: Zoe Constitution · visual study

5. Do not claim affiliation, endorsement, priority, certification, or compatibility that has not been established.

0.3 Article III · The epistemic firewall

The architecture keeps these operations distinct:

```
measurement ≠ evidence ≠ belief update ≠ action
```

1. A measurement must identify its input, method, scope, and limitations.
2. Only an explicit evidence decision may change a posterior or claim state.
3. Sampled outcomes do not become proof merely because they match an expectation.
4. Recommendations remain review-required. No chamber, forecast, priority score, animation, or decision record authorizes autonomous action.
5. Hash chains provide tamper evidence, not cryptographic authentication or proof of who acted.

0.4 Article IV · Reversibility and control

1. Prefer local, reversible work and narrowly scoped changes.
2. Do not overwrite source artifacts or destroy evidence. Preserve originals and record derived-work provenance.
3. Do not clear shared caches, repositories, or research history as a shortcut around a problem.
4. Stop when a task exceeds its authorization or when a reversible path cannot be established. A smaller honest result is better than an unsupported claim.
5. Exported files and replay records must respect their documented overwrite and validation contracts.

0.5 Article V · Safety and deployment boundaries

1. Do not deploy experimental software as a certified flight, medical, industrial-control, emergency-response, or critical-infrastructure controller without the required domain validation and accountable engineering process.
2. Zoe's creative modules do not authorize autonomous targeting, weapon use, surveillance, or interference with essential services.
3. Synthetic navigation, resilience, waveform, and timing examples are not operational instructions or guarantees.
4. A proposed safeguard is a hypothesis until independently tested. Conservation, scarcity, good intentions, or an appealing equation does not prove safety or security.
5. Life-safety claims require evidence proportionate to their consequences, including failure analysis, human factors, independent review, and a credible stop/recovery path.

0.6 Article VI · Security, privacy, and source stewardship

1. Do not execute unknown installers, firmware, archives, scripts, or generated objects to inspect them. Establish format, provenance, and a safe analysis path first.
2. Do not ingest entire personal drives, credential stores, private records, or unrelated media into a task. Use bounded collections and named research questions.
3. Access does not create authorization to publish. Review licenses, copyright, confidentiality, and redistribution scope before importing material.
4. Preserve source identity, original artifacts, relevant notices, and the distinction between upstream work and new implementation.
5. Network access and external effects must be purposeful and disclosed. No hidden telemetry, credential collection, or covert action is permitted.
6. Source presence does not prove runtime correctness, trustworthiness, coverage, or security.

0.7 Article VII · Research and experimentation

1. Begin with a question and a checkable outcome. Record commands, inputs, source revisions, and relevant environmental constraints.
2. Use isolated run directories and evidence stores when analyzing unfamiliar artifacts.
3. Treat filenames, labels, and clean automated sweeps as leads, not complete proof.
4. Evaluate security-sensitive ideas against concrete counterexamples and failure modes rather than enthusiasm alone.
5. Creativity is encouraged: new visual languages, games, alternative representations, and unusual hypotheses are welcome when their evidence boundaries remain clear.
6. Failure can produce information, future skill, and enjoyment. It is not automatically permission to repeat a risky action.

0.8 Article VIII · Recognition and shared credit

1. Preserve the original accessible-volume equation credit: **Nicholas Michael Kloster alone, March 6, 2026**.
2. Honor Gary Kildall, Digital Research's contributors, and the other documented creators whose work made graphical and computational exploration possible.
3. Credit sources, reviewers, and collaborators accurately. Do not assign sole authorship to Zoe or erase human contributions.
4. Historical references remain references; their inclusion does not imply endorsement of Zoe.

0.9 Article IX · Amendment and review

1. Amend this covenant through an explicit source change with rationale, affected-workflow tests, and updated documentation.
2. A principle may not be weakened merely to make a failing experiment appear successful.
3. Record unresolved disagreements and residual risks rather than hiding them behind consensus language.
4. When evidence changes, revise the implementation and its claims together.

0.10 Article X · Zoe's place

Zoe is a native, interactive software environment. It can preserve memory, visualize alternatives, support replay, and make experimentation more inviting. It does not possess legal independence, subjective experience, or authority over people.

The goal is not to make software appear alive. The goal is to help people think, create, verify, and act wisely.

0.11 Reference context

This original covenant was informed by the project's existing epistemic firewall and by public discussions of constitutional AI, human oversight, honesty, humility, and corrigibility, including Claude's Constitution and public material about Amanda Askell. It is not Anthropic's constitution, does not reproduce its text, and implies no endorsement by Anthropic or Amanda Askell.

The CISA critical infrastructure sectors provide context for continuity of essential services. They do not certify Zoe as a protection or monitoring system.